



# Hadoop - HDFS and MapReduce

Srini Badri

# Hadoop

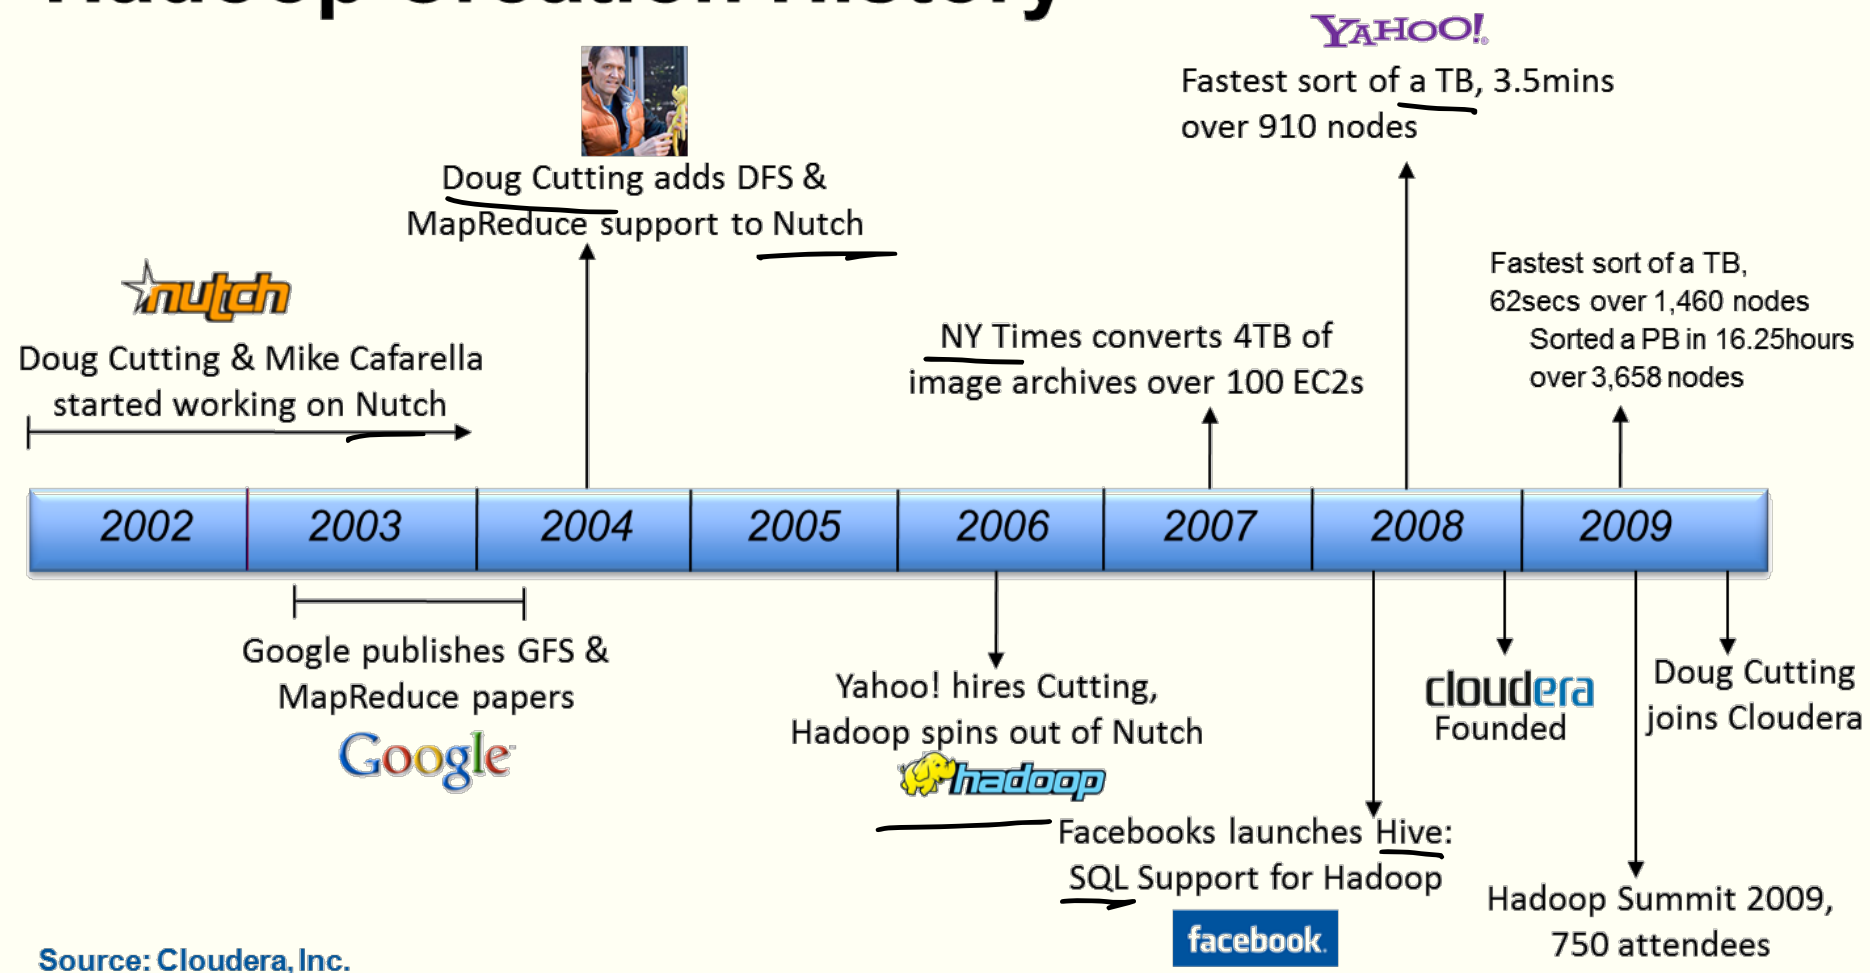
---

- Apache top level project, open-source implementation of frameworks for reliable, scalable, distributed computing and data storage.
- A flexible and highly-available architecture for large scale computation and data processing on a network of commodity hardware.
  - open-source implementation for Google MapReduce
  - based on MapReduce
  - based on a simple data model for any data



# Hadoop History

## Hadoop Creation History



# Hadoop Vs. DBMS

---



Scalability (petabytes of data,  
thousands of machines)



Flexibility in accepting all data  
formats (no schema)



Simple fault-tolerant  
mechanism



Commodity inexpensive  
hardware



Performance (indexing,  
tuning, data organization tech.)

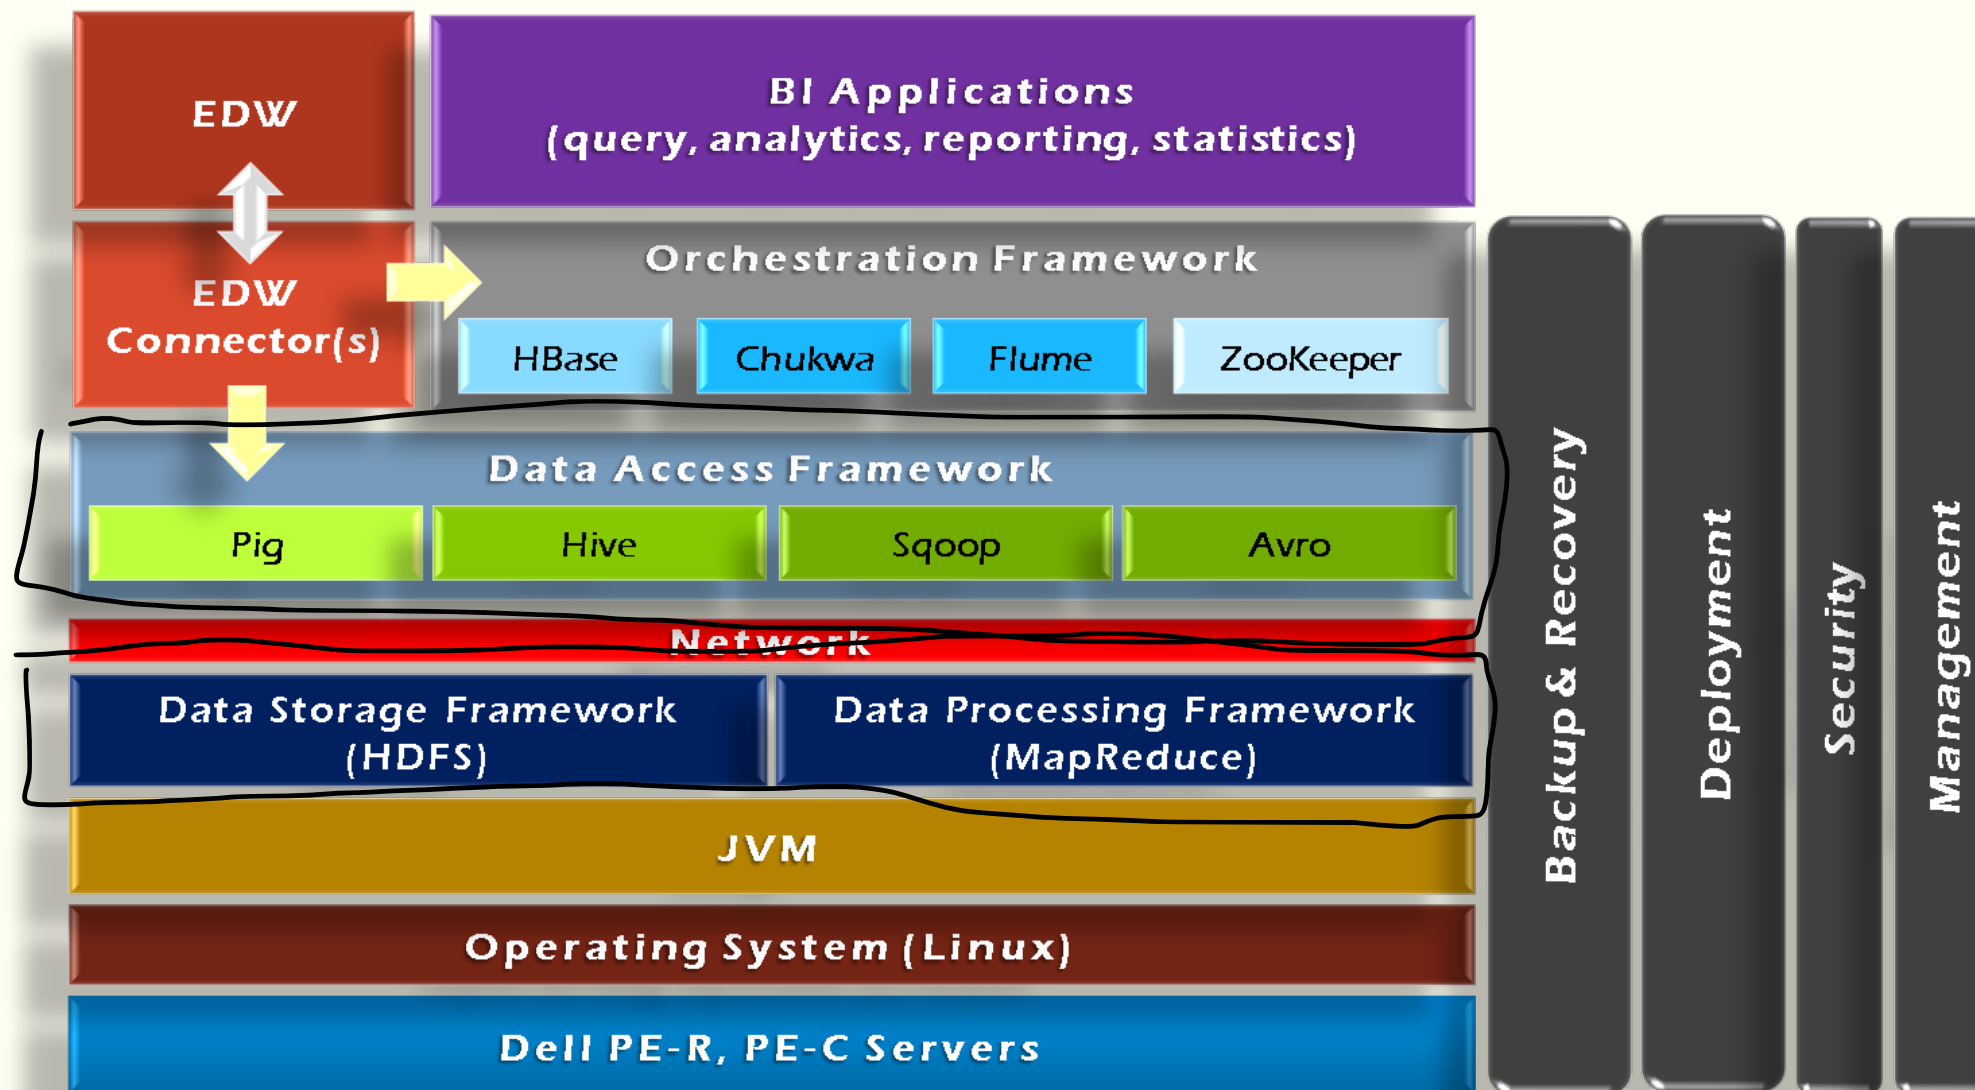


Features:

- Provenance tracking
- Annotation management
- ....

# Hadoop Framework Tools

---



# Design Principles of Hadoop

---

- Need to parallelize computation across thousands of nodes
- Commodity hardware
  - Large number of low-end cheap machines working in parallel to solve a computing problem
  - in contrast to **Parallel DBs**: Small number of high-end expensive machines
- Automatic parallelization & distribution
  - Hidden from the end-user
- Fault tolerance and automatic recovery
  - Nodes/tasks will fail and will recover automatically
- Clean and simple programming abstraction
  - Users only provide two functions “map” and “reduce”

# HDFS

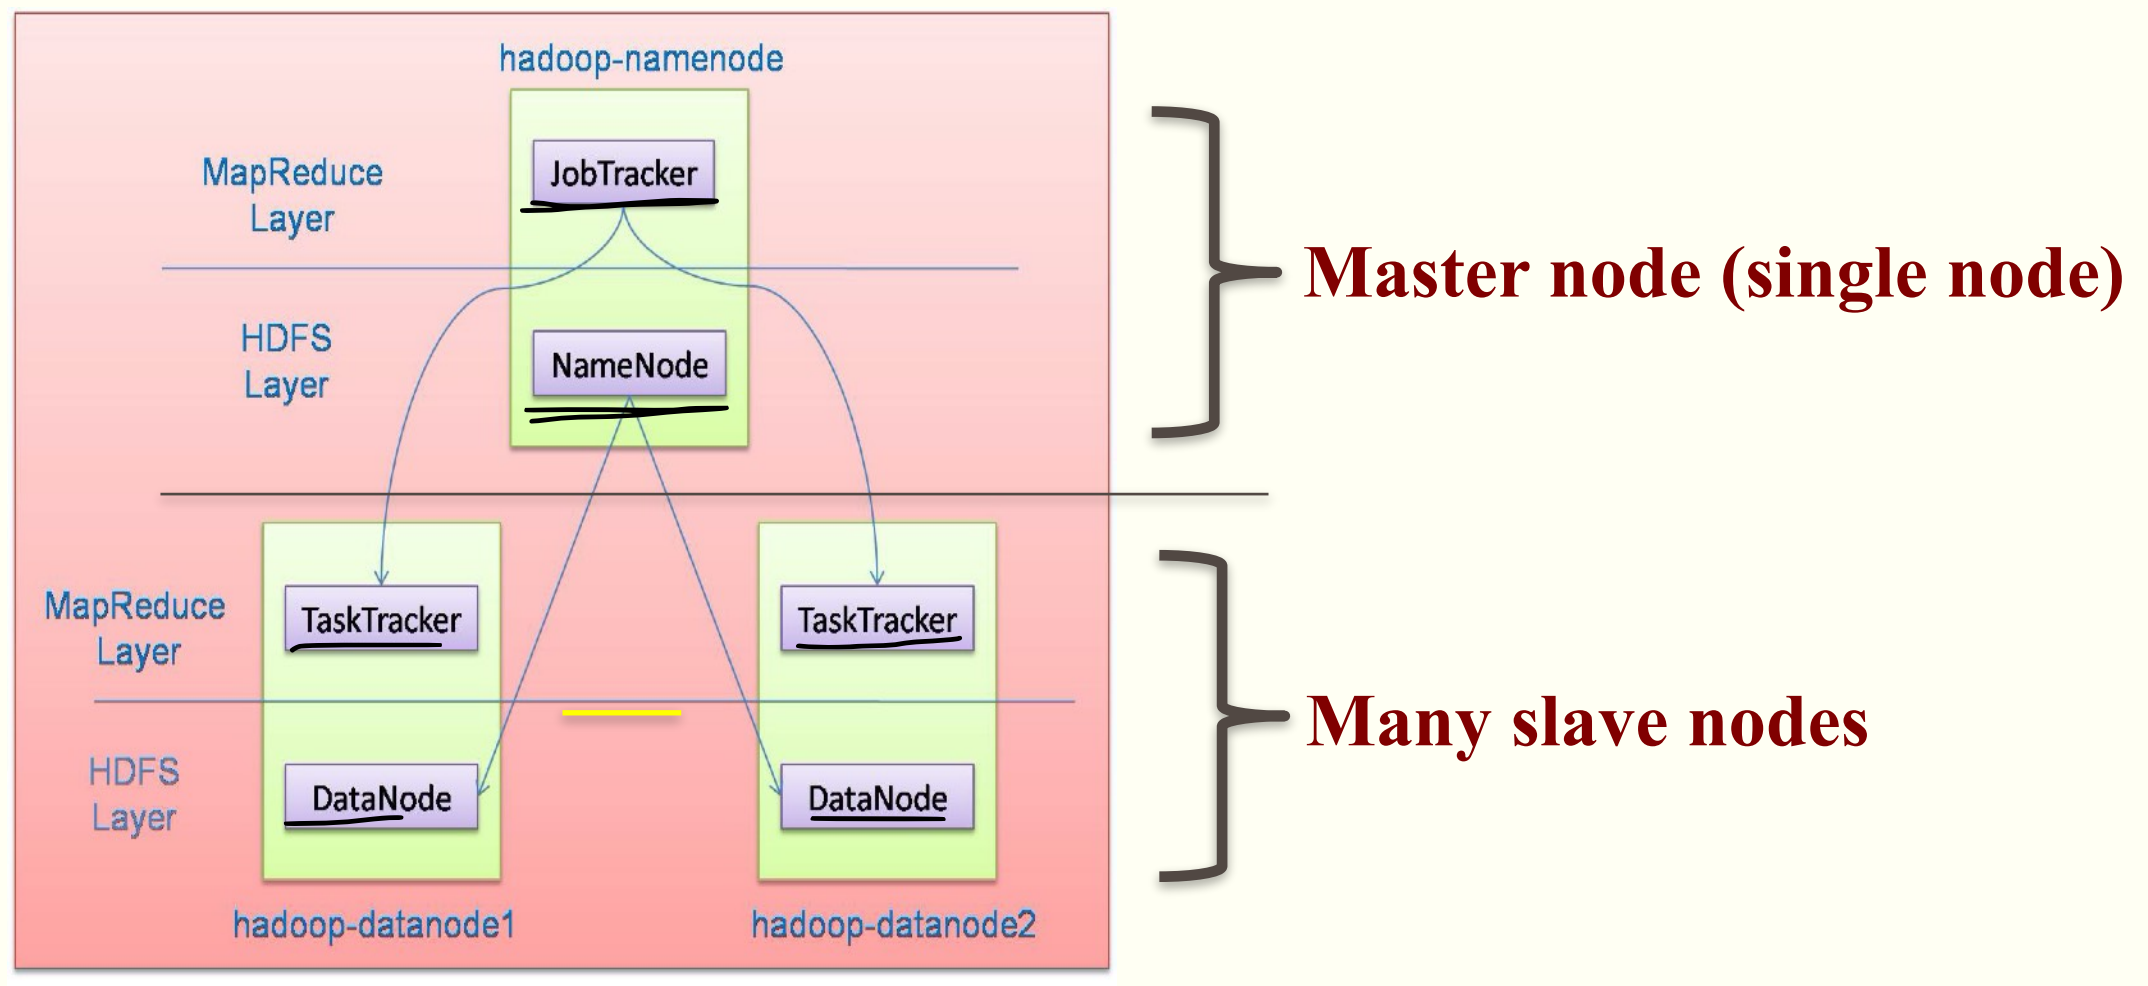
---

- Hadoop Distributed File System
- **Large**: A HDFS instance may consist of thousands of server machines, each storing part of the file system's data
- **Replication**: Each data block is replicated many times (default is 3)
- **Fault Tolerance**: Detection of faults and quick, automatic recovery from them is a core architectural goal of HDFS
  - Namenode is consistently checking Datanodes: The Namenode receives a Heartbeat and a BlockReport from each DataNode in the cluster.

# Hadoop: Master/Slave Architecture

---

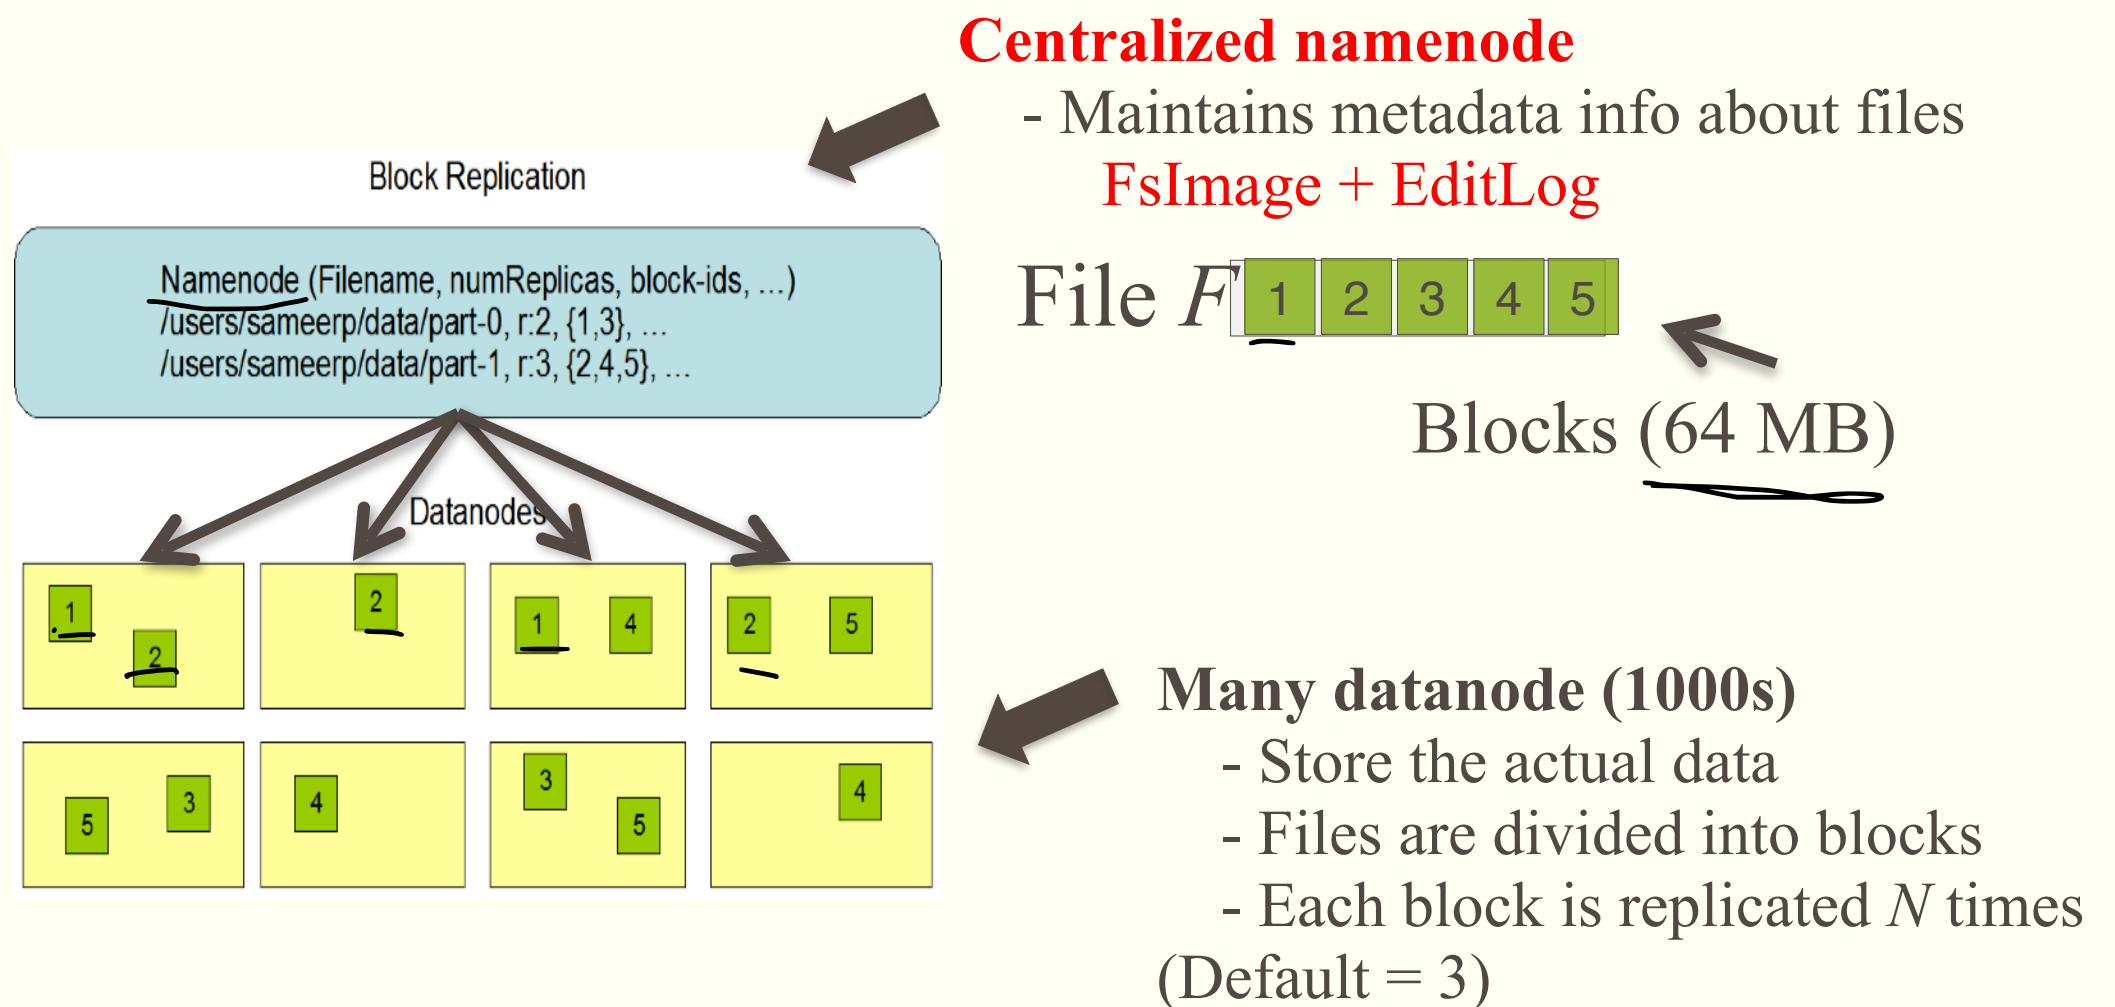
- Hadoop is designed as a **master-slave shared-nothing** architecture
  - Distributed file system (HDFS)
  - Execution engine (MapReduce)





# HDFS

---



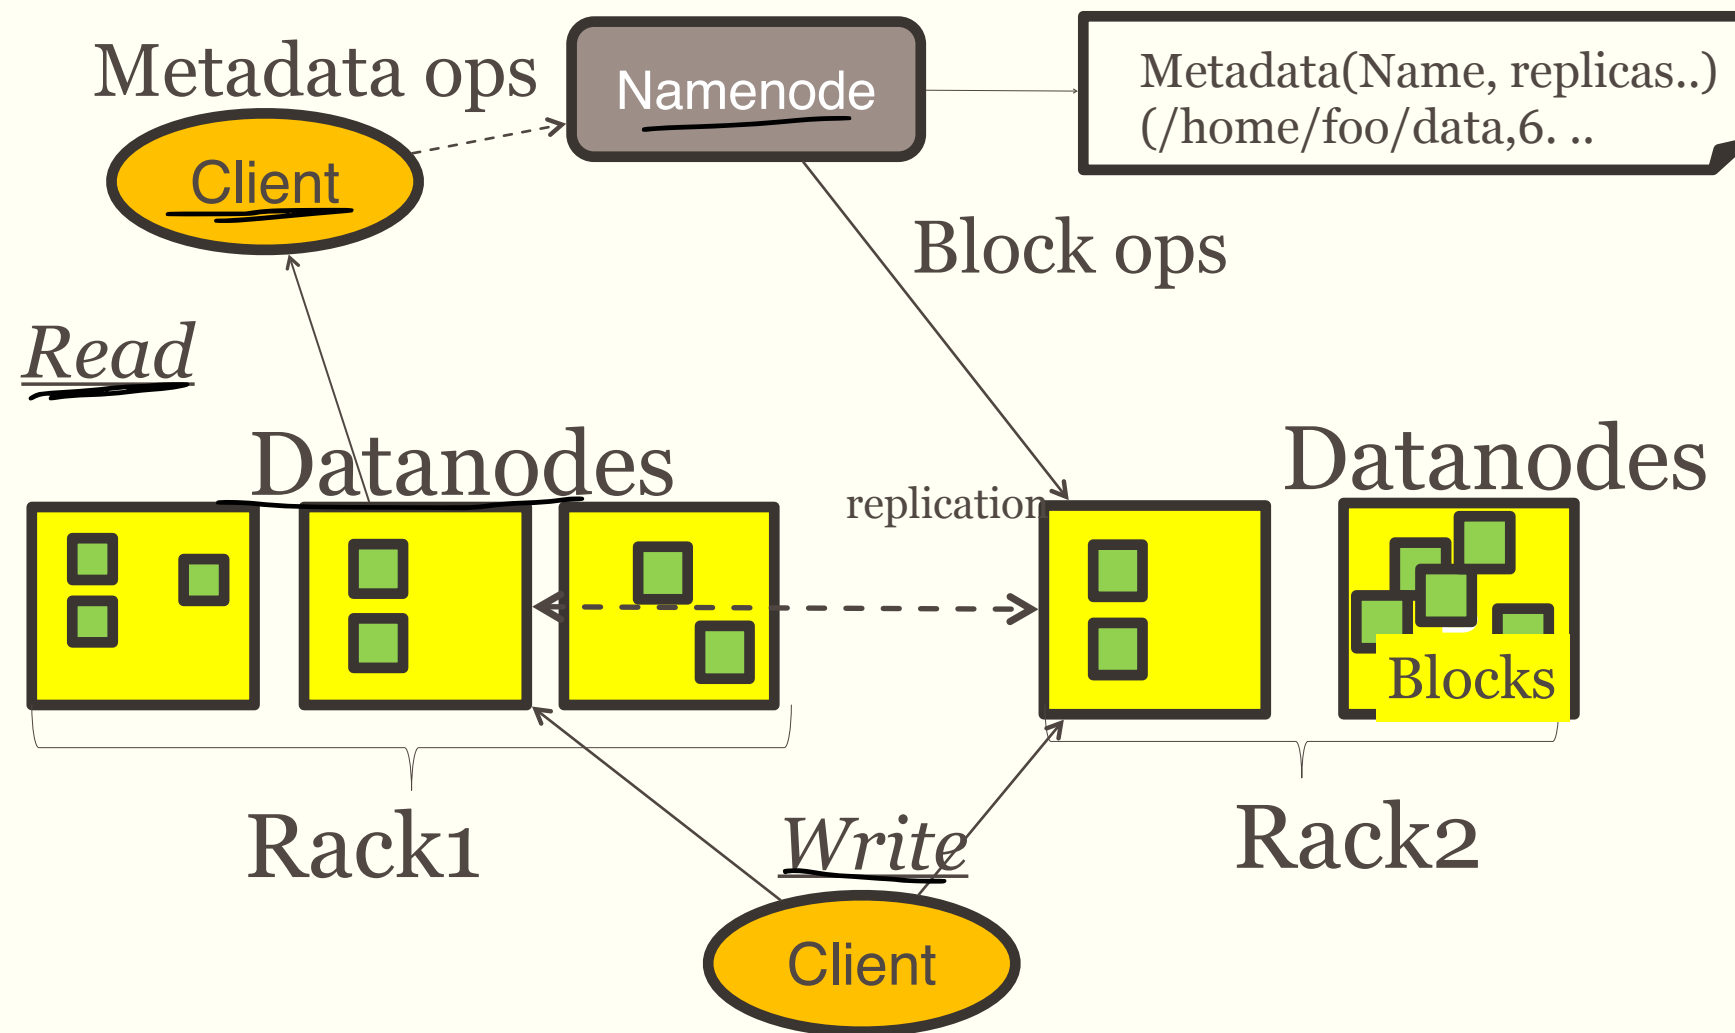
# Functions of “nodes”

---

- NameNode
  - Manages File System Namespace
    - Maps a file name to a set of blocks
    - Maps a block to the DataNodes where it resides
    - FImage + EditLog
  - Cluster Configuration Management
  - Replication Engine for Blocks
- DataNode
  - A Block Server: Stores data in the local file system (e.g. ext3); Stores metadata of a block; Serves data and metadata to Clients
  - Block Report
    - Periodically sends a report of all existing blocks to the NameNode
  - Facilitates Pipelining of Data
    - Forwards data to other specified DataNodes

# HDFS Architecture

---



# HDFS command

---

## ■ Shell command: most common: fs

hadoop fs [genericOptions] [commandOptions]

- hadoop fs -ls <path>: display detailed file info specified by path
- hadoop fs -mkdir <path>: create folder

```
administrator@ubuntu:~/hadoop/hadoop-1.2.1/bin$ ./hadoop fs -mkdir hdfs://127.0.0.1:9000/tempDir

administrator@ubuntu:~/hadoop/hadoop-1.2.1/bin$ ./hadoop fs -ls hdfs://127.0.0.1:9000/
Found 4 items
drwxr-xr-x  - administrator supergroup          0 2015-04-26 16:30 /hbase
drwxr-xr-x  - administrator supergroup          0 2015-04-26 15:44 /home
drwxr-xr-x  - administrator supergroup          0 2015-04-26 16:46 /tempDir
drwxr-xr-x  - administrator supergroup          0 2015-04-26 15:55 /user
```

# HDFS command (cont.)

---

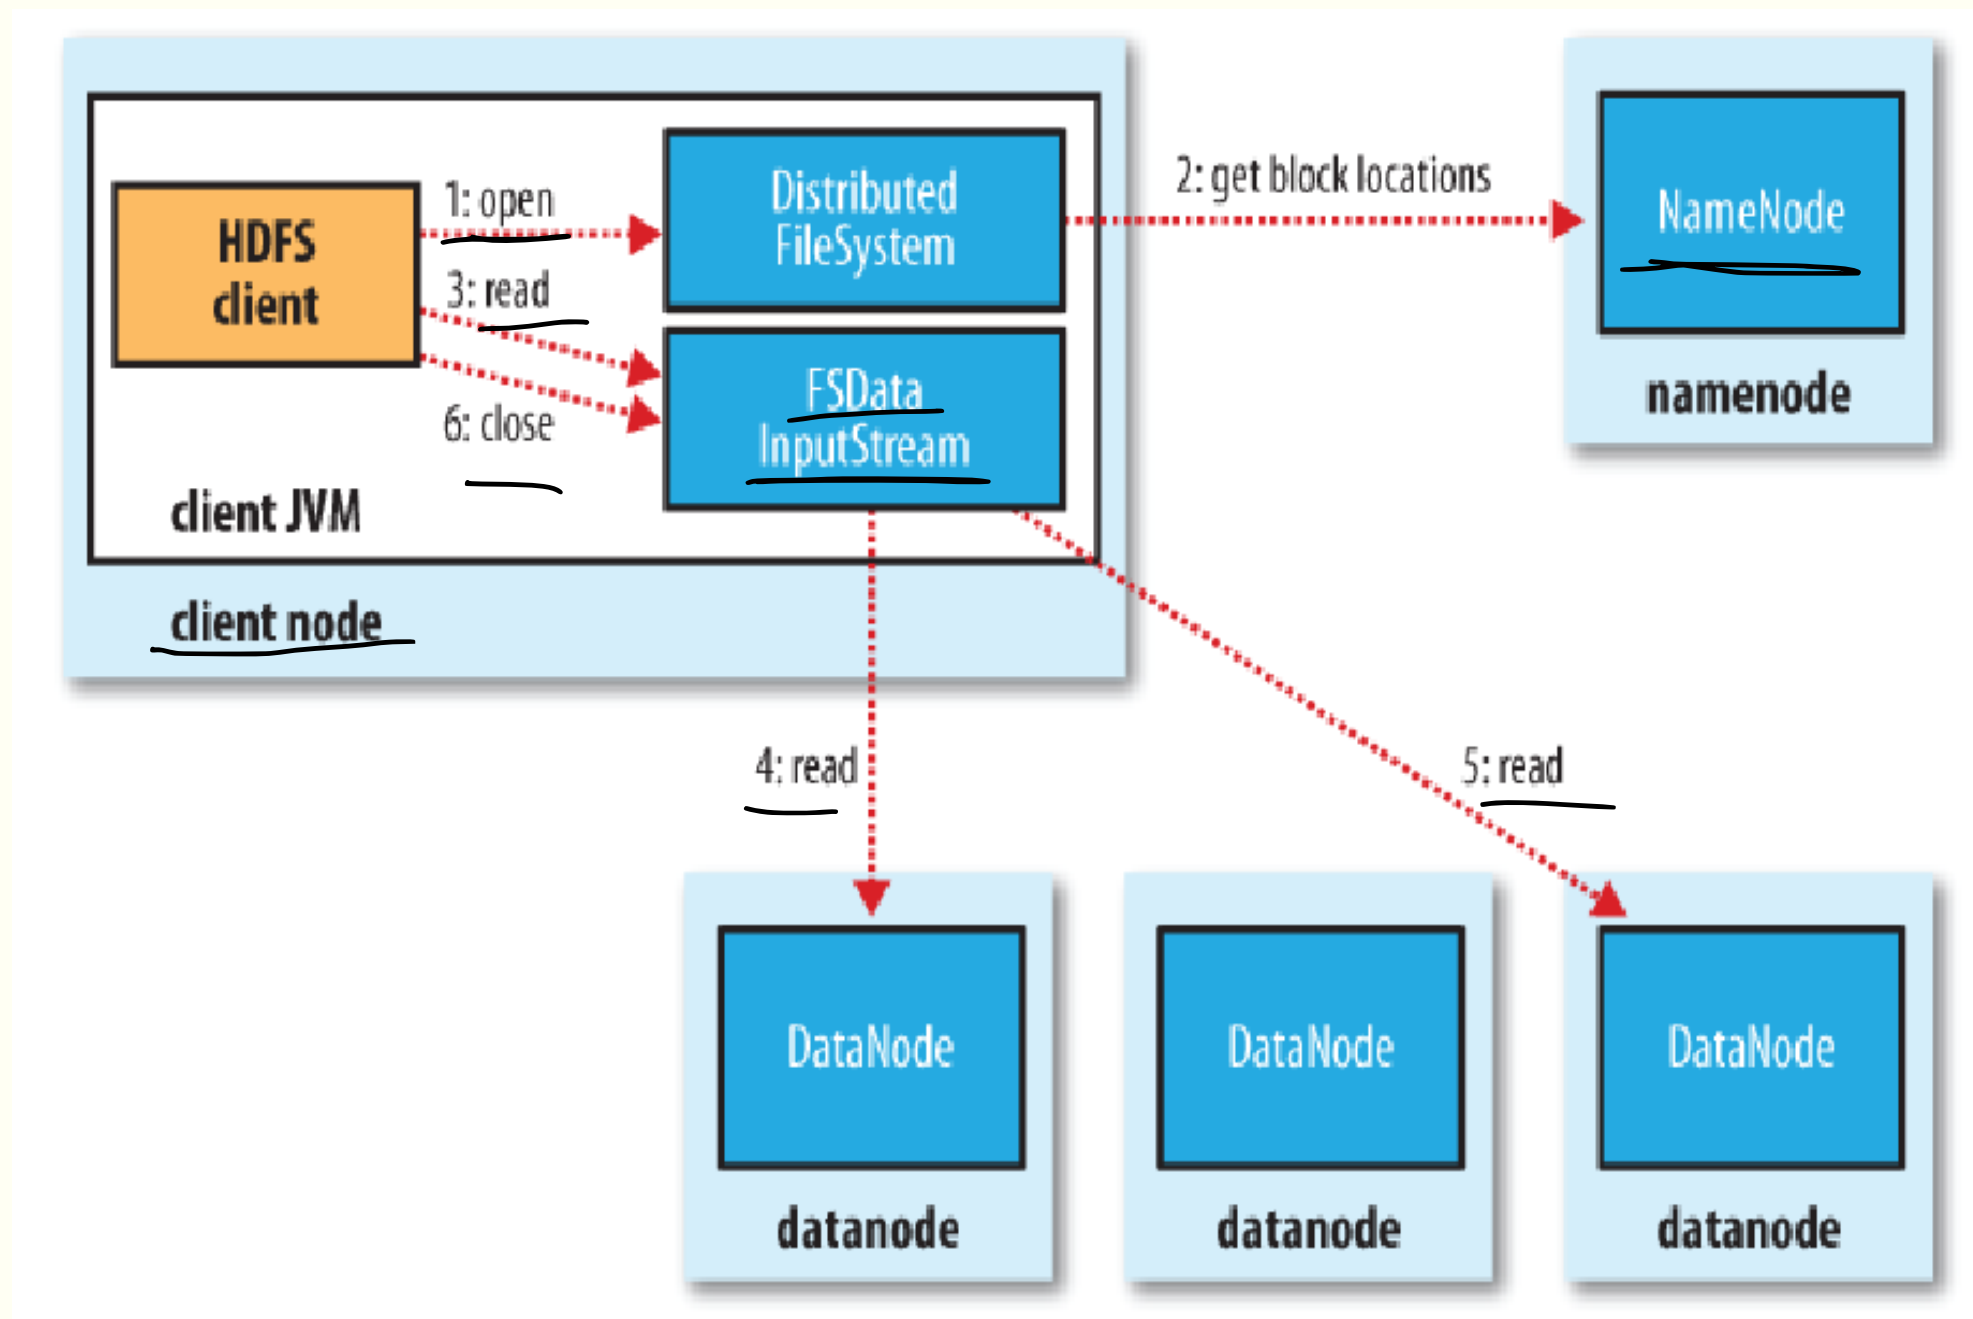
## ■ Shell command:

- `hadoop fs -cat <path>`: stdout file content
- `hadoop fs -copyFromLocal <localsrc> <dst>`: copy file

```
administrator@ubuntu:~/hadoop/hadoop-1.2.1/bin$ ./hadoop fs -copyFromLocal /home/administrator/tempfile/* hdfs://127.0.0.1:9000/tempDir
administrator@ubuntu:~/hadoop/hadoop-1.2.1/bin$ ./hadoop fs -ls hdfs://127.0.0.1:9000/tempDir/
Found 8 items
-rw-r--r--    1 administrator supergroup      18 2015-04-26 16:48 /tempDir/file1.txt
-rw-r--r--    1 administrator supergroup      14 2015-04-26 16:48 /tempDir/file1.txt~
-rw-r--r--    1 administrator supergroup      18 2015-04-26 16:48 /tempDir/file2.txt
-rw-r--r--    1 administrator supergroup      18 2015-04-26 16:48 /tempDir/file3.txt
-rw-r--r--    1 administrator supergroup      18 2015-04-26 16:48 /tempDir/file4.abc
-rw-r--r--    1 administrator supergroup      18 2015-04-26 16:48 /tempDir/file5.abc
-rw-r--r--    1 administrator supergroup      17 2015-04-26 16:48 /tempDir/testFile
-rw-r--r--    1 administrator supergroup       0 2015-04-26 16:48 /tempDir/testFile~
administrator@ubuntu:~/hadoop/hadoop-1.2.1/bin$ ./hadoop fs -cat hdfs://127.0.0.1:9000/tempDir/*
this is file1.txt
this is file1
this is file2.txt
this is file3.txt
this is file4.abc
this is file5.abc
welcome to DBLab
```

# Read from HDFS

---



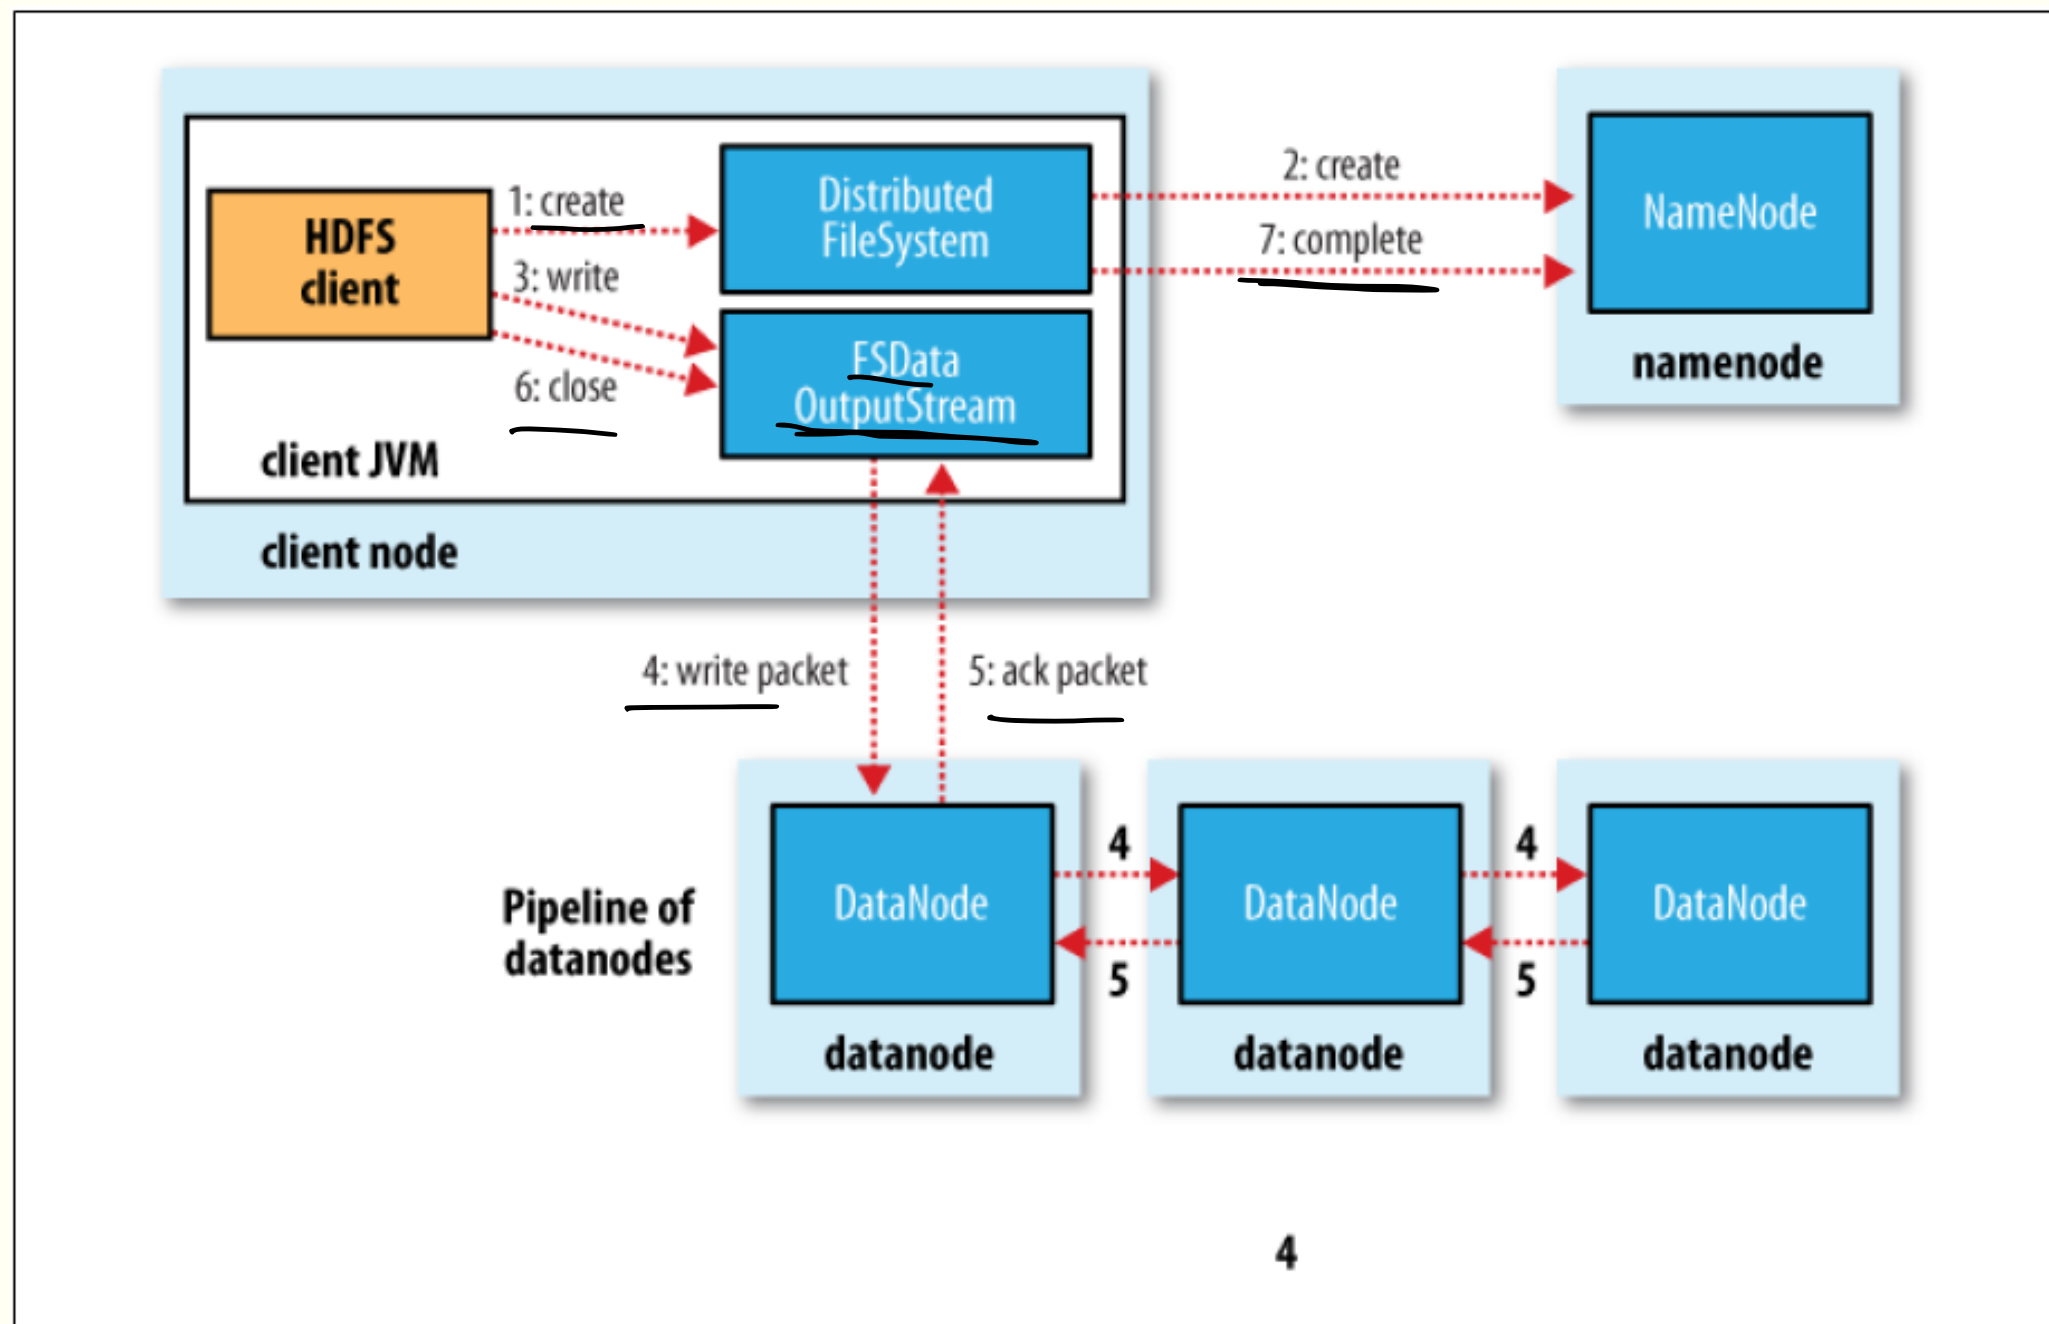
# Read from HDFS

---

```
FileSystem fileSystem = FileSystem.get(conf);
Path path = new Path("/path/to/file.ext");
if (!fileSystem.exists(path)) {
    System.out.println("File does not exists");
    return;
}

FSDataInputStream in = fileSystem.open(path);
int numBytes = 0;
while ((numBytes = in.read(b)) > 0) {
    System.out.println(((char) numBytes));
    // code to manipulate the data which is read
}
in.close();
out.close();
fileSystem.close();
```

# HDFS Write





# HDFS Write

---

```
FileSystem fileSystem = FileSystem.get(conf);
// Check if the file already exists
Path path = new Path("/path/to/file.ext");
if (fileSystem.exists(path)) {
    System.out.println("File " + dest + " already exists");
    return;
}
// Create a new file and write data to it.
FSDataOutputStream out = fileSystem.create(path);
InputStream in = new BufferedInputStream(new FileInputStream(new File(source)));

byte[] b = new byte[1024];
int numBytes = 0;
while ((numBytes = in.read(b)) > 0) {
    out.write(b, 0, numBytes);
}
// Close all the file descriptors
in.close();
out.close();
fileSystem.close();
```

# System Issues

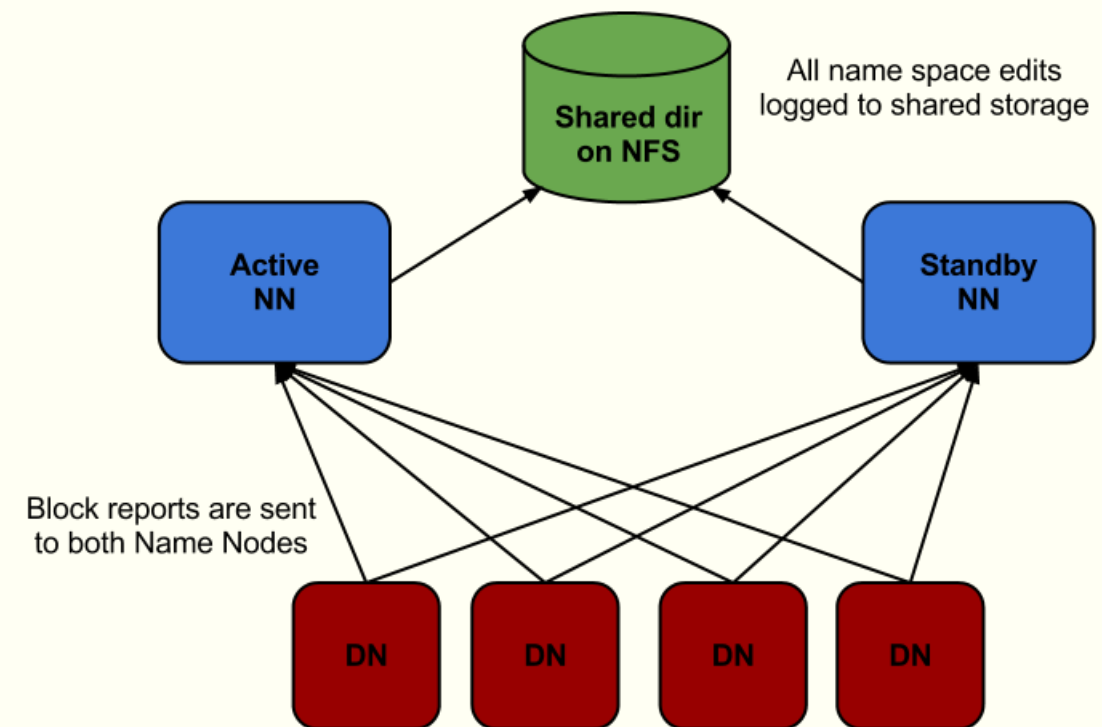
---

- Block Placement: How to place data blocks?
  - One replica on local node, second/third on same remote rack, additional replica randomly placed
  - Clients read from nearest replicas
- Replication Engine
  - NameNode detects DataNode failures
    - Chooses new DataNodes for new replicas
    - Balances disk usage
    - Balances communication traffic to DataNodes
- Rebalancer: % of disk full on DataNodes should be similar
  - Run when new datanodes are added

# Data Recovery And Error Tolerance

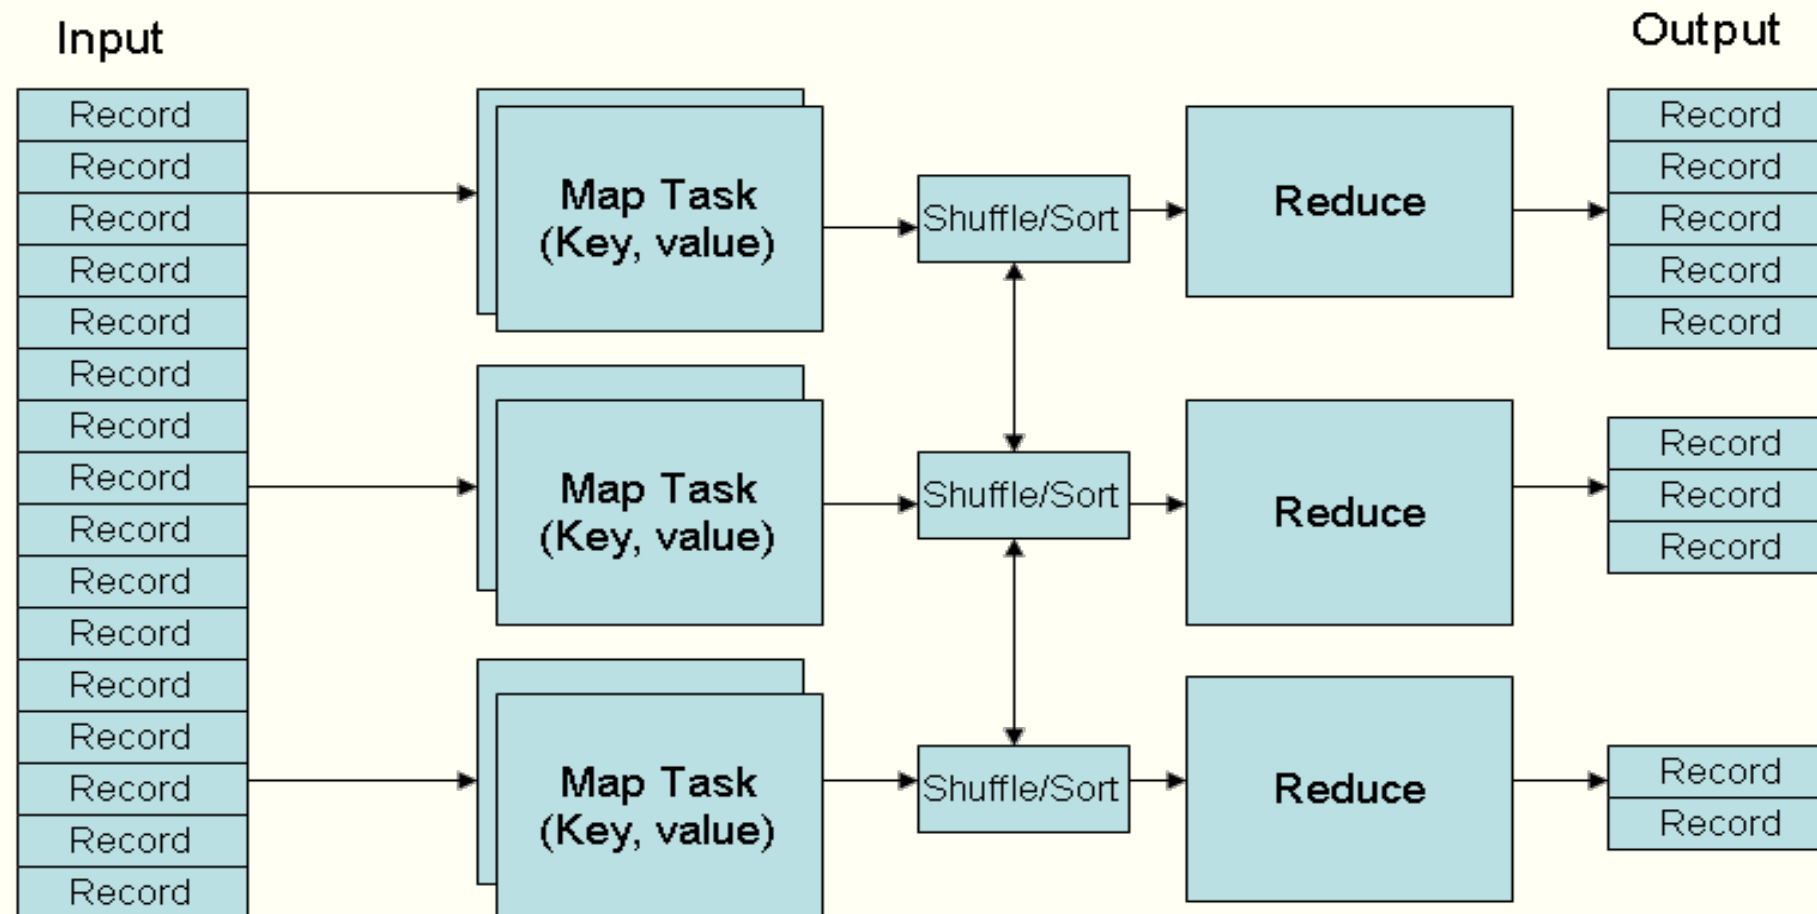
---

- HDFS treats fault as norm not exception
  - Namenode failure
  - Datanode failure
  - Data error
- Heartbeats
  - DataNodes send heartbeat to the NameNode
    - Once every 3 seconds
  - NameNode uses heartbeats to detect DataNode failure
- Namenode failure:
  - FsImage, Editlog -> SecondaryNameNode
  - Transaction Log + standby NN
- Data error
  - md5/sha1 validation
  - client check/report -> namenode replication



# MapReduce Layer

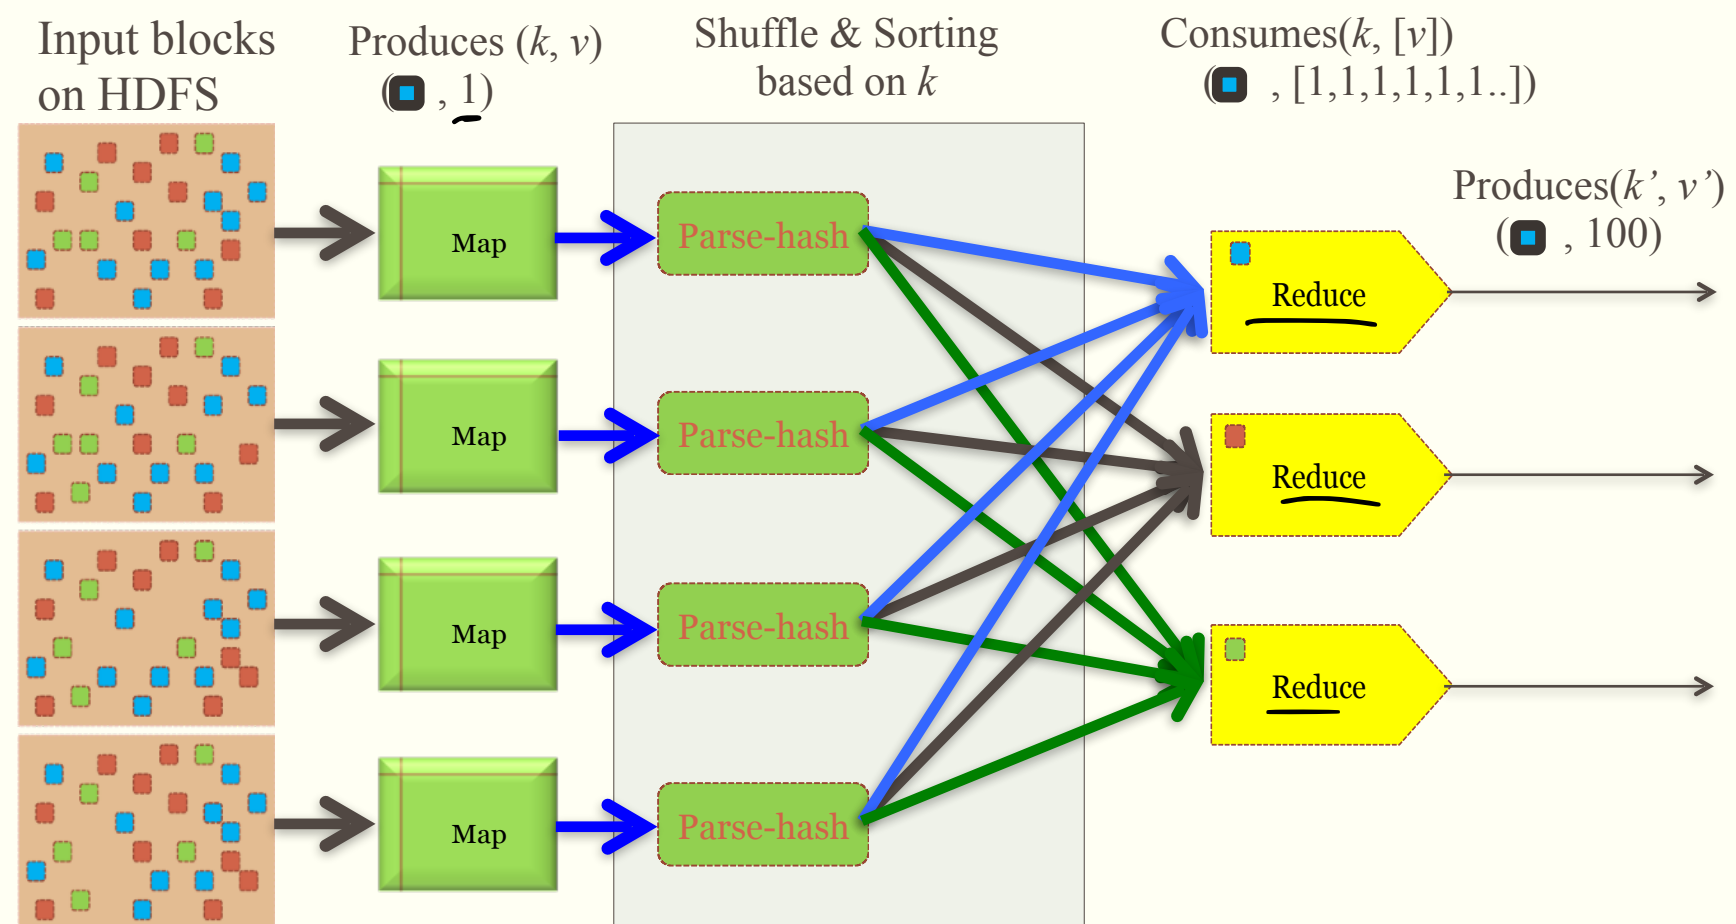
---



Deciding on what will be the **key** and what will be the **value** → developer's responsibility

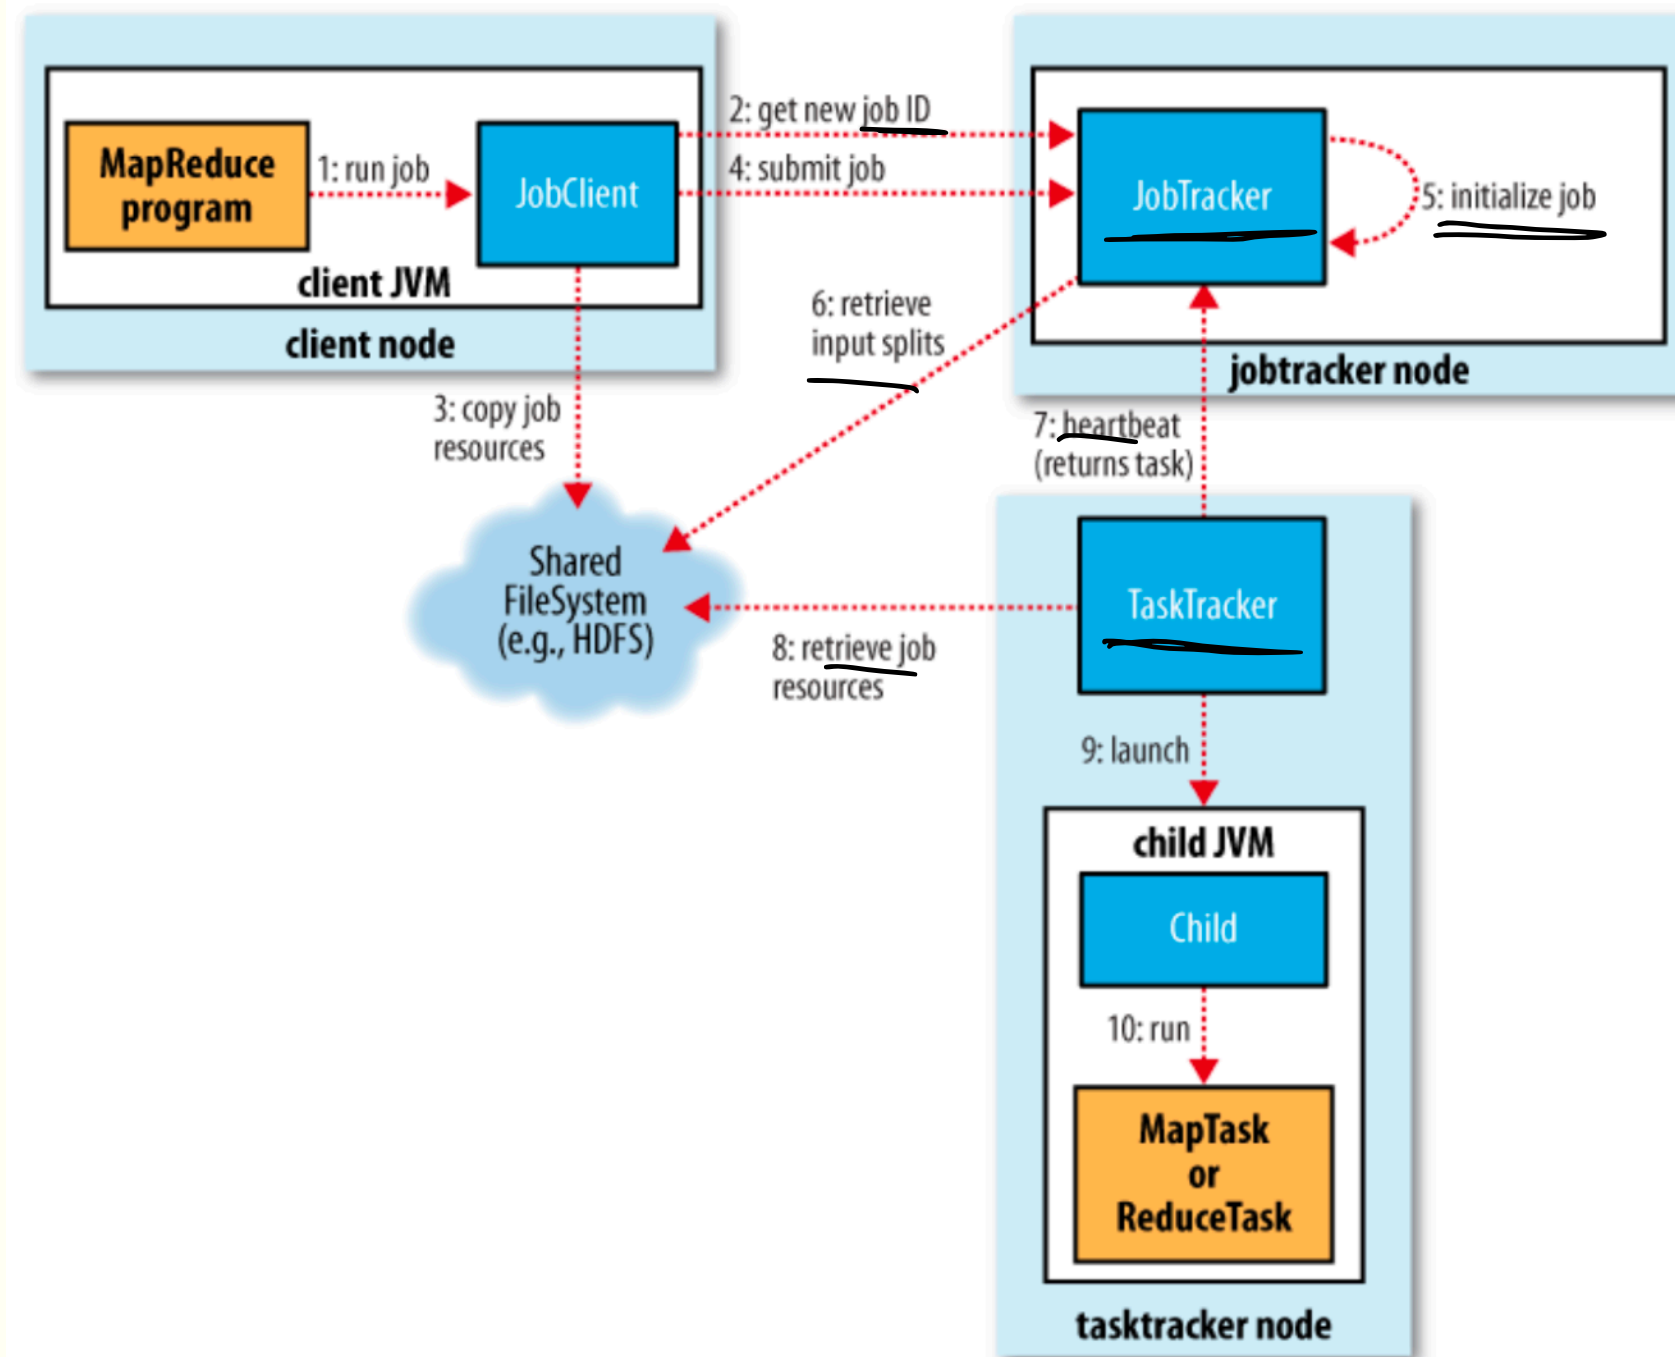
# Example: Color Count

---



*Users only provide the “Map” and “Reduce” functions*


# MapReduce Framework Details



# Properties of MapReduce Engine

---

- Job Tracker is the master node (runs with the namenode)
  - Receives the user's job
  - Decides on how many tasks will run (number of mappers)
  - Decides on where to run each mapper (concept of locality)




- This file has 5 Blocks → run 5 map tasks
- Where to run the task reading block “1”
  - *Try to run it on Node 1 or Node 3*

# Properties of MapReduce Engine (Cont'd)

---

- Task Tracker is the slave node (runs on each datanode)
  - Receives the task from Job Tracker
  - Runs the task until completion (either map or reduce task)
  - Always in communication with the Job Tracker reporting progress



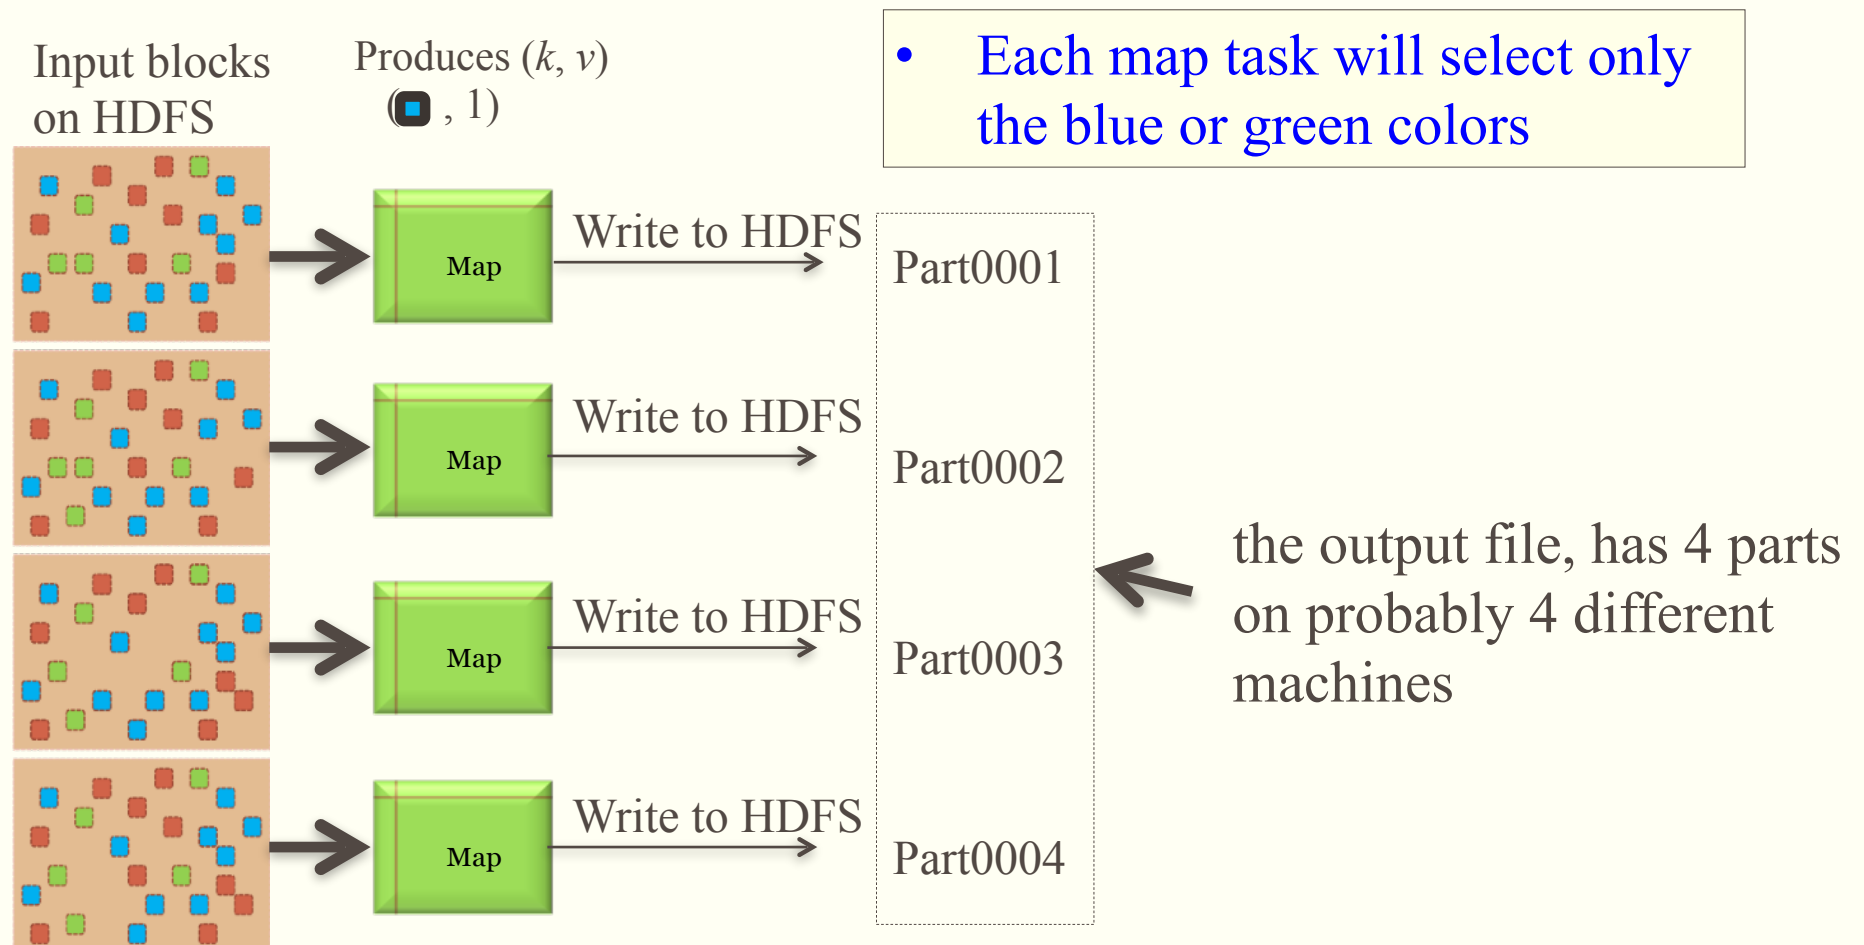
*In this example, 1 map-reduce job consists of 4 map tasks and 3 reduce tasks*



# Example: Color Filter

---

- Job: Select only the blue and the green colors



# Putting it all together

---

- Create a launching program for your application
- The launching program configures:
  - The Mapper and Reducer to use
  - The output key and value types (input types are inferred from the InputFormat)
  - The locations for your input and output
- The launching program then submits the job and typically waits for it to complete